

# Adversarial Search

Stephen G. Ware

CSCI 4525 / 5525



THE UNIVERSITY *of*  
NEW ORLEANS


# Search So Far

<b>Observable:</b>	Fully
<b>Agents:</b>	Single
<b>Deterministic:</b>	Deterministic
<b>Episodic:</b>	Sequential
<b>Static:</b>	Static
<b>Discrete:</b>	Discrete



# Adversarial Search

<b>Observable:</b>	Fully
<b>Agents:</b>	Multi
<b>Deterministic:</b>	Deterministic
<b>Episodic:</b>	Sequential
<b>Static:</b>	Static
<b>Discrete:</b>	Discrete




# Lessons to Learn

- Dealing with other agents
- Coping with large search spaces
- Making the best of limited resources
- Designing heuristics




# Game Theory Games

- Two perfectly rational players
- Each maximizes their own score
- Players take turns
- Moves are deterministic
- Game has a zero sum
- Perfect information is available to both players



# Tic-Tac-Toe Game Tree



# Utility

A **utility** function measures how desirable a given state is for an agent.

For Tic-Tac-Toe:

- 1 point if you won the game
- -1 point if you lost the game
- 0 points if the game is a tie



# Utility in 2PZS Games

Generally, each agent has its own utility function which it tries to maximize.


In the case of 2-player, zero-sum games, we can use a single utility function for both agents.

We name the two players:


- “max” tries to maximize the utility value
- “min” tries to minimize the utility value




# Tic-Tac-Toe Game Tree




# Game Trees in General




# Game Trees in General



# Game Trees in General



# Max's Best Move



# Minimax Search

To find the optimal strategy:

- Expand the tree using depth first search.
- Leaf nodes have a value equal to their utility.
- The value of a node at a max layer has the highest value of any of its children.
- The value of a node at a min layer has the lowest value of any of its children.

# Minimax Search

```
function min_max(Node n) {  
    best = find_max(n);  
    return the move that results in best utility;  
}
```



# Recursive Depth First Tree Search

```
function dfs(Node n) {  
    for every child of n {  
        child = next child of n;  
        dfs(child);  
    }  
}
```





# Find Max

```
function find_max(Node n) {  
    if n is a leaf node, return utility(n);  
    best =  $-\infty$ ;  
    for every child of n {  
        child = next child of n;  
        child_value = find_min(child);  
        best = max(best, child_value);  
    }  
    return best;  
}
```

# Find Min

```
function find_min(Node n) {  
    if n is a leaf node, return utility(n);  
    best =  $+\infty$ ;  
    for every child of n {  
        child = next child of n;  
        child_value = find_max(child);  
        best = min(best, child_value);  
    }  
    return best;  
}
```

# Minimax Example

max's turn



min's turn



# Minimax Example

max's turn

$-\infty$

min's turn

$+\infty$

# Minimax Example

max's turn

$-\infty$

min's turn

$+\infty$

3

# Minimax Example

max's turn

$-\infty$

min's turn

3


3

# Minimax Example

max's turn



min's turn



# Minimax Example

max's turn

$-\infty$

min's turn

-5

3

-5



# Minimax Example

max's turn



min's turn



# Minimax Example

max's turn

-5

min's turn

-5

3

-5


7

# Minimax Example

max's turn




min's turn



# Minimax Example

max's turn


min's turn



# Minimax Example

max's turn


min's turn



# Minimax Example

max's turn


min's turn



# Minimax Example

max's turn


min's turn



# Minimax Example


max's turn

min's turn







# Minimax Example




# Minimax Example




# Minimax Example




# Minimax Example




# Minimax Example



# Minimax Example



# Minimax Example



# Handling Large Trees

What if we don't have time to expand the entire game tree?

A estimate of the Chess game tree states that it has about  $35^{100}$  nodes.

That's over

25000  
000  
000  
00000000000000000000 nodes.



# Alpha Beta Pruning


- We don't need to know the utility of each branch of the game tree, only the good moves.
- If we find a branch which is worse than the best outcome we have found so far, we don't need to know how much worse. Just ignore it.
- Don't bother considering a move if we know it will be worse than one we've already found.

# Alpha Beta Pruning

- Keep track of two numbers during search:
- $\alpha$ , the highest utility found so far
- $\beta$ , the lowest utility found so far

At max levels, don't bother with values  $\geq \beta$ .

At min levels, don't bother with values  $\leq \alpha$ .



# Alpha Beta Pruning


- Keep track of two numbers during search:
- $\alpha$ , the highest utility found so far
- $\beta$ , the lowest utility found so far

At max levels, don't bother with values  $\geq \beta$ .

(because min won't let that happen)

At min levels, don't bother with values  $\leq \alpha$ .

(because max won't let that happen)




# Minimax Search + ABP

```
function min_max_ab(Node n) {  
    best = find_max_ab(n,  $-\infty$ ,  $+\infty$ );  
    return the move that creates best utility;  
}
```




# Find Max + ABP

```
function find_max_ab(Node n,  $\alpha$ ,  $\beta$ ) {  
    if n is a leaf node, return utility(n);  
    best =  $-\infty$ ;  
    for every child of n {  
        child = next child of n;  
        child_value = find_min_ab(child,  $\alpha$ ,  $\beta$ );  
        best = max(best, child_value);  
        if best  $\geq \beta$ , return best;  
         $\alpha$  = max( $\alpha$ , best);  
    }  
    return best;  
}
```




# Find Min + ABP

```
function find_min_ab(Node n,  $\alpha$ ,  $\beta$ ) {  
    if n is a leaf node, return utility(n);  
    best =  $+\infty$ ;  
    for every child of n {  
        child = next child of n;  
        child_value = find_max_ab(child,  $\alpha$ ,  $\beta$ );  
        best = min(best, child_value);  
        if best  $\leq \alpha$ , return best;  
         $\beta$  = min( $\beta$ , best);  
    }  
    return best;  
}
```



# Alpha Prune Example

max's turn



$$\alpha = -\infty$$

$$\beta = +\infty$$

min's turn

# Alpha Prune Example

max's turn

$$\alpha = -\infty$$

$$\beta = +\infty$$

min's turn

$$\alpha = -\infty$$

$$\beta = +\infty$$



# Alpha Prune Example

max's turn

$$\alpha = -\infty$$

$$\beta = +\infty$$

min's turn

$$\alpha = -\infty$$

$$\beta = +\infty$$

3

# Alpha Prune Example

max's turn

$$\alpha = -\infty$$

$$\beta = +\infty$$

min's turn

$$\alpha = -\infty$$

$$\beta = 3$$

3


# Alpha Prune Example

max's turn

$$\alpha = -\infty$$
$$\beta = +\infty$$

min's turn

$$\alpha = -\infty$$
$$\beta = 3$$



# Alpha Prune Example

max's turn


$$\alpha = -\infty$$

$$\beta = +\infty$$

min's turn

$$\alpha = -\infty$$

$$\beta = -1$$



# Alpha Prune Example

max's turn

$$\alpha = -\infty$$

$$\beta = +\infty$$

min's turn

$$\alpha = -\infty$$

$$\beta = -1$$



# Alpha Prune Example

max's turn

$$\alpha = -\infty$$

$$\beta = +\infty$$

min's turn

-1

$$\alpha = -\infty$$

$$\beta = -1$$

3

-1

4

# Alpha Prune Example

max's turn

$$\alpha = -1$$
$$\beta = +\infty$$

min's turn

-1

$$\alpha = -\infty$$
$$\beta = -1$$

3

-1

4

# Alpha Prune Example

max's turn

$$\alpha = -1$$
$$\beta = +\infty$$

min's turn

-1

$$\alpha = -\infty$$
$$\beta = -1$$

$$\alpha = -1$$
$$\beta = +\infty$$





# Alpha Prune Example

max's turn


$$\alpha = -1$$
$$\beta = +\infty$$

min's turn

-1

$$\alpha = -\infty$$
$$\beta = -1$$

$$\alpha = -1$$
$$\beta = +\infty$$



# Alpha Prune Example

max's turn


$$\alpha = -1$$
$$\beta = +\infty$$

min's turn

-1

$$\alpha = -\infty$$
$$\beta = -1$$

$$\alpha = -1$$
$$\beta = +\infty$$



# Alpha Prune Example

max's turn

$$\alpha = -1$$
$$\beta = +\infty$$

min's turn

-1

$$\alpha = -\infty$$
$$\beta = -1$$

$$\alpha = -1$$
$$\beta = +\infty$$



# Alpha Prune Example

max's turn

$$\alpha = -1$$
$$\beta = +\infty$$

min's turn

-1


$$\alpha = -\infty$$
$$\beta = -1$$

-2


$$\alpha = -1$$
$$\beta = +\infty$$




# Alpha Prune Example




# Alpha Prune Example



# Alpha Prune Example




# Alpha Prune Example







# Alpha Prune Example




# Alpha Prune Example



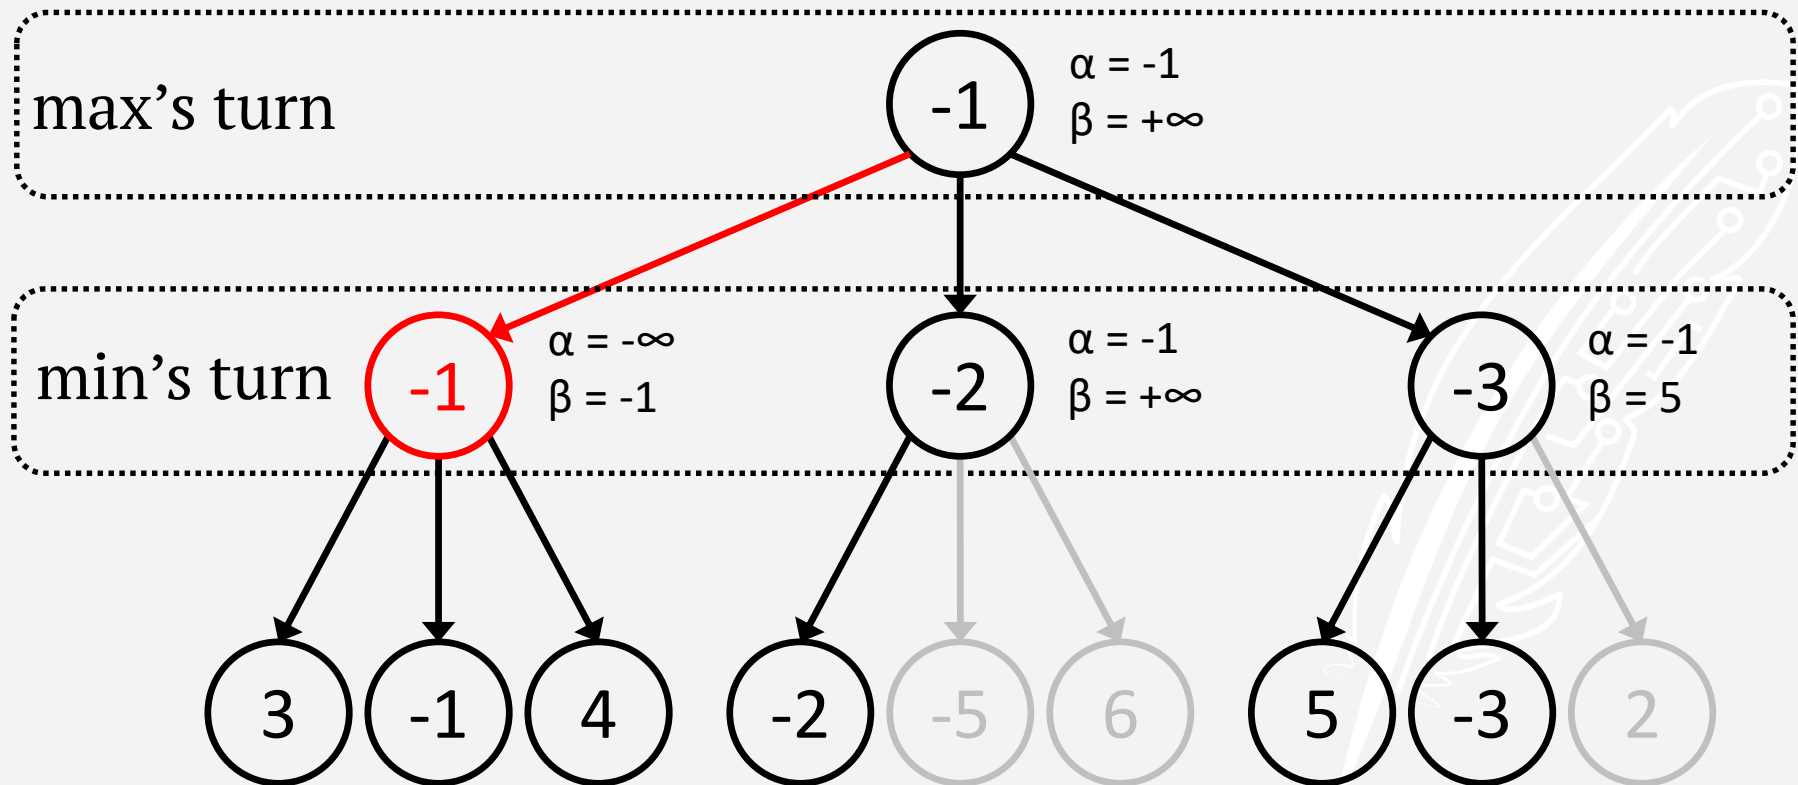
# Alpha Prune Example



# Alpha Prune Example



# Alpha Prune Example



# Alpha Beta Pruning

Returns exactly the same results as minimax, but can expand significantly fewer branches.




# Pruned Branches

We do not know the exact value of a branch which has been pruned... only that it is high enough or low enough that we shouldn't bother with it.




# Alpha Prune Example





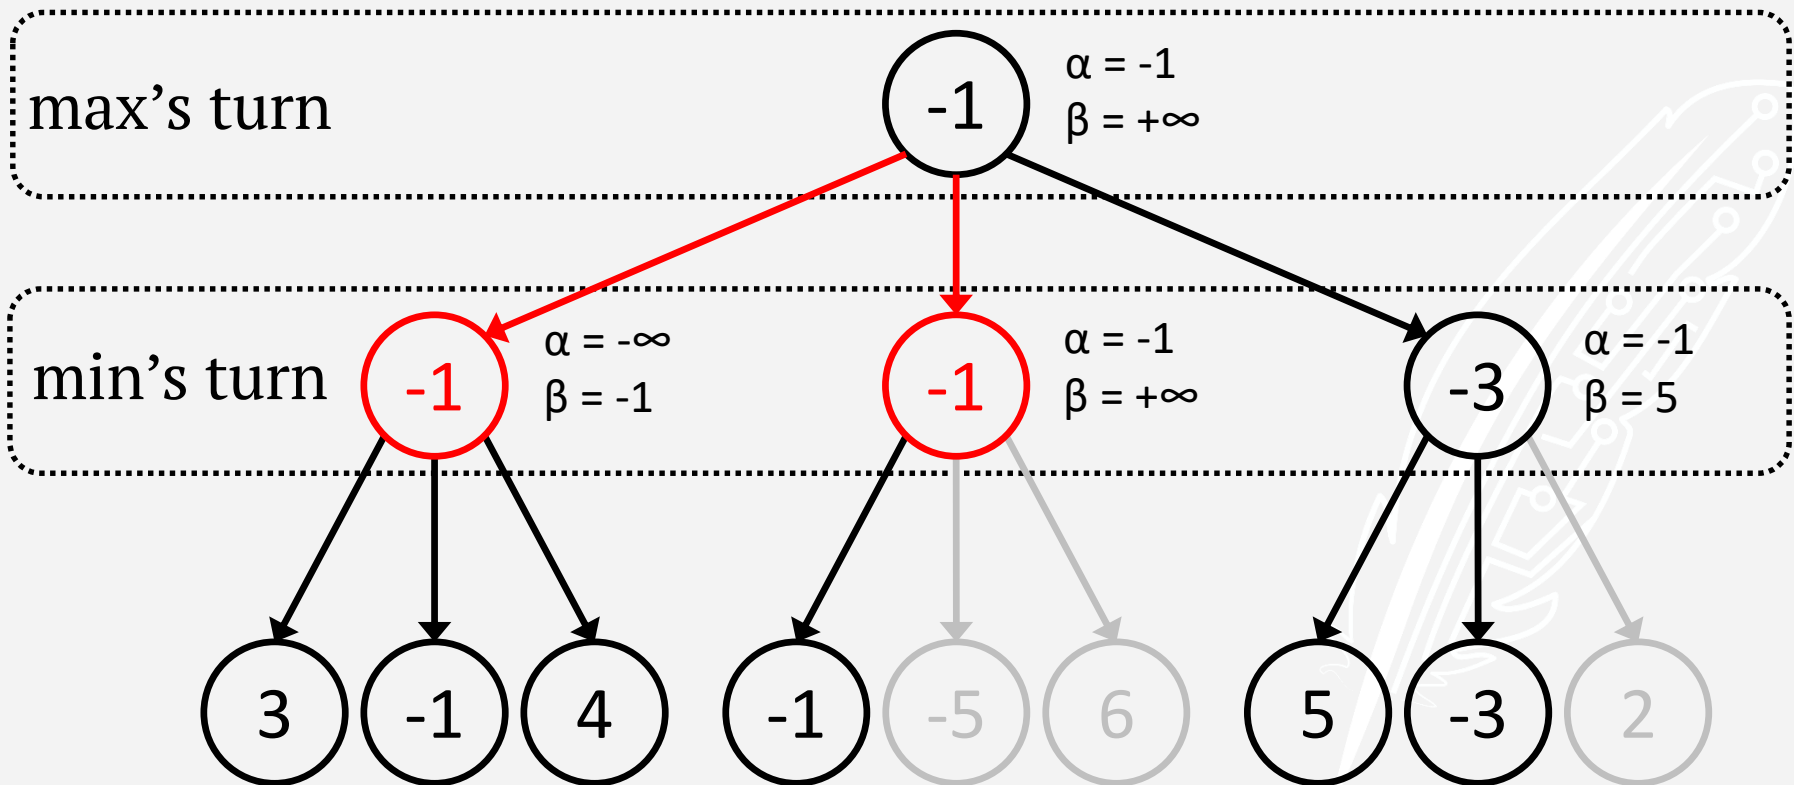
# Alpha Prune Example

Does it matter which one we choose?



# Alpha Prune Example

Does it matter which one we choose? Yes! Don't choose the pruned branch.



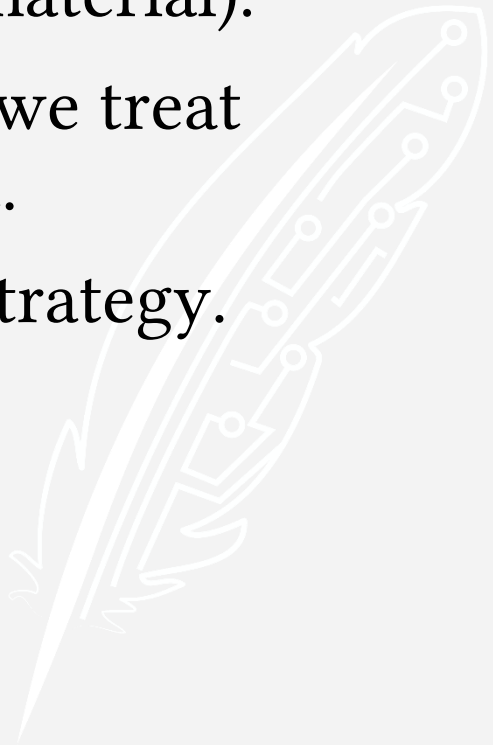
# Limited Time and Space

It is often impractical to go all the way to a leaf.




# Imperfect Decision Making

- Design a utility function which can evaluate the state of a game in progress (e.g. chess material).
- Impose an arbitrary cut off past which we treat nodes as leaf nodes even if they are not.
- No longer guaranteed to find the best strategy.



# Forward Pruning

- **Beam Search** limits the number of children that a given parent can have.
- Example: Consider only the 5 best children of a node, and don't bother with the rest.
- Dangerous because we might prune the best or worst branch.



# Iterative Deepening


- If we run out of time or memory before minimax search is finished, we can't trust the answer it gives us.
- One solution is to use iterative deepening.
- Chess example: First look 2 moves ahead. If that finishes, look 4 moves ahead. If that finishes, look 6 moves ahead. Keep increasing this value until we run out of time or space.

(Make sure not to trust the last, incomplete search!)

# Move Ordering

- The sooner we find very high and very low moves, the more nodes can be pruned.
- Consider moves in an order that is likely to discover high and low values sooner.
- Chess example: always consider moves which capture a piece first. These are likely to lead to more extreme utility values and thus improve pruning.

# Horizon Problem



## Material

White: 7

Black: 7


Black to move.

Black looks 2 moves ahead.

If the bishop is captured, black's score will drop to 4.



# Horizon Problem




## Material

White: 7

Black: 7



# Horizon Problem



## Material

White: 7


Black: 6

Black to move.

Black looks 2 moves ahead.

If the bishop is captured, black's score will drop to 3.

# Horizon Problem




**Material**

White: 7

Black: 6



# Horizon Problem



## Material

White: 7

Black: 5



# Quiescence

- When search is arbitrarily cut off, if the next move is really good or really bad, we won't see it.
- Only cut off the search if we have reached the max depth limit AND the game is in a relative state of equilibrium (e.g. next moves are not big ones).
- Chess example: Cut off search if we have reached the max depth AND the previous move was not a capture.